

Chapter 3 Design Theory for Relational Databases

Contents

- Functional Dependencies
- Decompositions
- Normal Forms (BCNF, 3NF)
- Multivalued Dependencies (and 4NF)
- Reasoning About FD's + MVD's

Remember our questions:

- Why do we design relations like the example? -- **good design**
- What makes a good relational database schema? -- **no redundancy, no Update/delete anomalies,**
- what we can do if it has flaws? -- **decomposition**

New Question:

- **any standards for a good design?**
→ **Normal forms: a condition on a relation schema that will eliminate problems**
- **any standards or methods for a decomposition?**

Boyce-Codd Normal Form

- We say a relation R is in **BCNF** if whenever $X \rightarrow Y$ is a nontrivial FD that holds in R , X is a superkey.
 - Remember: *nontrivial* means Y is not contained in X .
 - Remember, a *superkey* is any superset of a key (not necessarily a proper superset).

Example

Drinkers(name, addr, beersLiked, manf, favBeer)

FD's: name->addr favBeer, beersLiked->manf

- Only key is {name, beersLiked}.
- In each FD, the left side is *not* a superkey.
- Any one of these FD's shows *Drinkers* is not in BCNF

Another Example

Beers(name, manf, manfAddr)

FD's: name->manf, manf->manfAddr

- Only key is {name} .
- name->manf does not violate BCNF, but manf->manfAddr does.


Decomposition into BCNF

- Given: relation R with FD's F .
- Look among the **given FD's** for a BCNF violation $X \rightarrow Y$.
 - If any FD following from F violates BCNF, then there will surely be an FD in F itself that violates BCNF.
- Compute X^+ .
 - Not all attributes, or else X is a superkey.

Decompose R Using $X \rightarrow Y$

- Replace R by relations with schemas:
 1. $R_1 = X^+$.
 2. $R_2 = R - (X^+ - X)$.
- *Project* given FD's F onto the two new relations.

Decomposition Picture



Example: BCNF Decomposition

Drinkers(name, addr, beersLiked, manf, favBeer)

$F = \text{name} \rightarrow \text{addr}, \quad \text{name} \rightarrow \text{favBeer}, \quad \text{beersLiked} \rightarrow \text{manf}$

- Pick BCNF violation $\text{name} \rightarrow \text{addr}$.
- Close the left side: $\{\text{name}\}^+ = \{\text{name}, \text{addr}, \text{favBeer}\}$.
- Decomposed relations:
 1. Drinkers1(name, addr, favBeer)
 2. Drinkers2(name, beersLiked, manf)

Example -- Continued

- We are not done; we need to check Drinkers1 and Drinkers2 for BCNF.
- Projecting FD's is easy here.
- For **Drinkers1**(name, addr, favBeer), relevant FD's are **name->addr** and **name->favBeer**.
 - Thus, {**name**} is the only key and Drinkers1 is in BCNF.

Example -- Continued

- For $\text{Drinkers2}(\underline{\text{name}}, \underline{\text{beersLiked}}, \text{manf})$, the only FD is $\text{beersLiked} \rightarrow \text{manf}$, and the only key is $\{\text{name}, \text{beersLiked}\}$.
 - Violation of BCNF.
- $\text{beersLiked}^+ = \{\text{beersLiked}, \text{manf}\}$, so we decompose *Drinkers2* into:
 1. $\text{Drinkers3}(\underline{\text{beersLiked}}, \text{manf})$
 2. $\text{Drinkers4}(\underline{\text{name}}, \underline{\text{beersLiked}})$

Example -- Concluded

- The resulting decomposition of *Drinkers* :
 1. *Drinkers1*(name, addr, favBeer)
 2. *Drinkers3*(beersLiked, manf)
 3. *Drinkers4*(name, beersLiked)
- Notice: *Drinkers1* tells us about drinkers, *Drinkers3* tells us about beers, and *Drinkers4* tells us the relationship between drinkers and the beers they like.

Classroom Exercise

- Any two-attribute relation $R(A,B)$ is in BCNF

Several Cases:

- 1) No nontrivial FD's
- 2) $A \rightarrow B$
- 3) $A \rightarrow B, B \rightarrow A$

Third Normal Form -- Motivation

- There is one structure of FD's that causes trouble when we decompose.
- $R(A,B,C)$ $AB \rightarrow C$ and $C \rightarrow B$.
 - Example: A = street address, B = city, C = zip code.
- There are two keys, $\{A,B\}$ and $\{A,C\}$.
- $C \rightarrow B$ is a BCNF violation, so we must decompose into $R_1(AC)$, $R_2(BC)$.

We Cannot Enforce FD's

- The problem is that if we use AC and BC as our database schema, we cannot enforce the FD $AB \rightarrow C$ by checking FD's in these decomposed relations.
- Example with $A = \text{street}$, $B = \text{city}$, and $C = \text{zip}$ on the next slide.

A = street, B = city,
and C = zip

An Unenforceable FD

Zip → city
keeps

street	zip
545 Tech Sq.	02138
545 Tech Sq.	02139

city	zip
Cambridge	02138
Cambridge	02139

Join tuples with equal zip codes.

street	city	zip
545 Tech Sq.	Cambridge	02138
545 Tech Sq.	Cambridge	02139

Although no FD's were violated in the decomposed relations,
FD **street city -> zip** is **violated** by the database as a whole.

3NF Let's Us Avoid This Problem

- 3rd Normal Form (3NF) modifies the BCNF condition so we do not have to decompose in this problem situation.
- An attribute is *prime* if it is a member of any key.
- $X \rightarrow A$ violates 3NF if and only if X is not a superkey, and also A is not prime.

Example: 3NF

- In our problem situation $R(A,B,C)$ with FD's $AB \rightarrow C$ and $C \rightarrow B$, we have keys AB and AC .
- Thus A , B , and C are each prime.
- Although $C \rightarrow B$ violates BCNF, it does not violate 3NF.

BCNF vs. 3NF

	conditions	example
BCNF	If $X \rightarrow Y$ is a nontrivial FD that holds in R , X is a superkey.	$R(A,B,C)$ with $A \rightarrow B$, $A \rightarrow C$
3NF	If $X \rightarrow Y$ is a nontrivial FD that holds in R , X is a superkey, or Y is a prime	$R(A,B,C)$ with $AB \rightarrow C$ and $C \rightarrow B$.

2 NF: no nonkey attribute is dependent on only a portion of the primary key. $R(A,B,C)$ with $A \rightarrow B$, $B \rightarrow C$

1 NF: every component of every tuple is an atomic value.

Properties of a decomposition

- Elimination of anomalies by a decomposition, it needs other two properties:
 1. *Lossless Join* : it should be possible to project the original relations onto the decomposed schema, and then reconstruct the original.
 2. *Dependency Preservation* : it should be possible to check in the projected relations whether all the given FD's are satisfied.

What 3NF and BCNF Give You

- We can get (1: *Lossless Join*) with a **BCNF decomposition**.
- We can get both (1) and (2: *Dependency Preservation*) with **a 3NF decomposition**.
- But we can't always get (1) and (2) with a BCNF decomposition.
 - street-city-zip is an example.

Testing for a Lossless Join

- If we project R onto R_1, R_2, \dots, R_k , can we recover R by rejoining?
- Any tuple in R can be recovered from its projected fragments.
- So the only question is: when we rejoin, do we ever get back something we didn't have originally?

Example

- Any tuple in R can be recovered from its projected fragments.
- Not less, **not more**.

As long as FD $b \rightarrow c$ holds, the joining of two projected tuples cannot produce a bogus tuple

A	B	C
1	2	3
4	2	5



A	B	B	C
1	2	2	3
4	2	2	5




A	B	C
1	2	3
1	2	5
4	2	3
4	2	5

The Chase Test (an example)

To prove that a tuple t in the join, using FD's in F , also to be a tuple in R .

- Suppose tuple t comes back in the join.
- Then t is the join of projections of some tuples of R , one for each R_i of the decomposition.
- Can we use the given FD's to show that one of these tuples must be t ?



The Chase Test – (method)

- Start by assuming $t = abc\dots$
- For each i , there is a tuple s_i of R that has a , b , c, \dots in the attributes of R_i .
- s_i can have any values in other attributes.
- We'll use the same letter as in t , but with a subscript, for these components.

Example: The Chase

- Let $R = ABCD$, and the decomposition be AB , BC , and CD .
- Let the given FD's be $C \rightarrow D$ and $B \rightarrow A$.
- Suppose the tuple $t = abcd$ is the **join** of tuples projected onto AB , BC , CD .

The tuples
of R pro-
jected onto
AB, BC, CD.

The *Tableau*

<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>
<i>a</i>	<i>b</i>	<i>c</i> ₁	<i>d</i> ₁
<i>a</i>₂ <i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>₂ <i>d</i>
<i>a</i> ₃	<i>b</i> ₃	<i>c</i>	<i>d</i>

Use *B* -> *A*

Use *C* -> *D*

We've proved the
second tuple must be *t*.

Summary of the Chase Test method

1. If two rows agree in the left side of a **FD**, make their right sides agree too.
2. Always **replace** a subscripted symbol by the corresponding unsubscripted one, if possible.
3. If we ever get an unsubscripted row, we know any tuple in the project-join is in the original (**the join is lossless**).
4. Otherwise, the final tableau is a **counterexample**.

Example: Lossy Join (more tuples)

- Same relation $R = ABCD$ and same decomposition: AB , BC , and CD .
- But with only the FD $C \rightarrow D$.

These projections
rejoin to form

abcd

The *Tableau*

A	B	C	D
<i>a</i>	<i>b</i>	<i>c</i> ₁	<i>d</i> ₁
<i>a</i> ₂	<i>b</i>	<i>c</i>	<i>d</i> ₂
<i>a</i> ₃	<i>b</i> ₃	<i>c</i>	<i>d</i>

These three tuples are an example
R that shows the join lossy. *abcd*
is not in *R*. but we can project and
rejoin to get *abcd*. **More tuples**

Use *C* > *D*

Some results (until now)

- Some decompositions can not keep lossless join (lossy join).
- Use **chase method** to find out whether the decomposition is lossy join.
- **BCNF decomposition is lossless join**, sometimes it can not keep functional dependencies.
- Relations with 3NF keep lossless join and also functional dependencies.
- How to decompose relations to reach 3NF?

3NF Synthesis Algorithm

- We can always construct a decomposition into **3NF relations** with a **lossless join** and **dependency preservation**.
- Need *minimal basis* for the FD's:
 1. Right sides are single attributes.
 2. No FD can be removed.
 3. No attribute can be removed from a left side.

Constructing a Minimal Basis

1. Split right sides.
2. Repeatedly try to remove an FD and see if the remaining FD's are equivalent to the original.
3. Repeatedly try to remove an attribute from a left side and see if the resulting FD's are equivalent to the original.

3NF Synthesis – method

1. Find a minimal basis for F
2. One relation for each FD in the minimal basis.
 1. Schema is the union of the **left and right sides**.
 2. $X \rightarrow A$ then (XA) is a schema.
3. If no key is contained in an FD, then **add one relation whose schema is some key**.

Algorithm 3.26 is on pp.103

Example: 3NF Synthesis

- Relation $R = ABCD$.
- FD's $A \rightarrow B$ and $A \rightarrow C$. Key is AD
- Decomposition: AB and AC from the FD's, plus AD for a key.

Another example

- Relation $R(A,B,C,D,E)$
- FD's $AB \rightarrow C, C \rightarrow B, A \rightarrow D$ key?

3NF synthesis:

- 1) A minimal basis
- 2) $R_1(ABC)$ $R_2(CB)$ $R_3(AD)$
- 3) R_2 is a part of R_1 , delete R_2
- 4) No key is in R_1, R_3 , add a key $R_4(ABE)$

R has two keys: ABE, ACE. Add one of them

Classroom exercises

Given $R(A,B,C,D,E)$ FD's $AB \rightarrow C$, $C \rightarrow B$, $A \rightarrow D$

- To test $R_1(ABC)$, $R_3(AD)$, $R_4(ABE)$ is in 3NF
- To test whether functional dependency keeps in the R_1, R_3, R_4
- To test the decomposition is lossless.

Why It Works (3NF synthesis)

- **Preserves dependencies**: each FD from a minimal basis is contained in a relation, thus preserved.
- **Lossless Join**: use the chase to show that the row for the relation that contains a key can be made all-unsubscripted variables.

Question:

Why we say “BCNF decomposition” , “3NF synthesis”?

BCNF decomposition algorithm

Input: relation R + FDs for R

Output: decomposition of R into BCNF relations

With “lossless join”

1. Compute keys for R

2. Repeat until all relations are in BCNF:

Pick any R' with $A \rightarrow B$ that violates BCNF

Decompose R' into $R_1(A^+)$ and $R_2(A, \text{rest})$

Compute FDs for R_1 and R_2

Compute keys for R_1 and R_2

3NF synthesis

- Input: relation R + FDs for R
- Output: decomposition of R into 3NF

With “lossless join” and keep FD’s

1. Computer key of R
2. Find a minimal basis for F
3. One relation for each FD in the minimal basis.
4. If no key is contained in an FD, then add one relation whose schema is some key.

Summary

- **Conditions** of Norm Forms (BCNF, 3NF)
- The **way to decompose** in order to reach BCNF
- The **way to decompose** in order to reach 3NF
- The **way to test** the join is lossless join