

Chapter 6 The database


Language SQL –as a **tutorial**

- **About SQL**

SQL is a standard database language, adopted by many commercial systems.


ANSI SQL, SQL-92 or SQL2, **SQL99 or SQL3 extends SQL2 with object-relational features. SQL2003 is the collection of extensions to SQL3.**

- **How to query the database**
- **How to make modifications on database**
- **Transactions in SQL**



Union, Intersection and Difference

- **Union, intersection, and difference of relations are expressed by the following forms, each involving subqueries:**
 - **(subquery) UNION (subquery)**
 - **(subquery) INTERSECT (subquery)**
 - **(subquery) EXCEPT (subquery)**




Example

- **From relations**

Likes(drinker, beer), Sells(bar, beer, price) and Frequents(drinker, bar), find the drinkers and beers such that:

- 1. The drinker likes the beer, and**
- 2. The drinker frequents at least one bar that sells the beer.**




Solution

**(SELECT * FROM Likes)
INTERSECT**

The drinker frequents
a bar that sells the
beer.

**(SELECT drinker, beer
FROM Sells, Frequents
WHERE Frequents.bar = Sells.bar
);**



Bag Semantics

- **Although the SELECT-FROM-WHERE statement uses bag semantics, the default for union, intersection, and difference is set semantics.**
 - **That is, duplicates are eliminated as the operation is applied.**




Motivation: Efficiency

- **When doing projection in relational algebra, it is easier to avoid eliminating duplicates.**
 - Just work tuple-at-a-time.
- **When doing intersection or difference, it is most efficient to sort the relations first.**
 - At that point you may as well eliminate the duplicates anyway.



Controlling Duplicate Elimination


- Force the result to be a set by **SELECT DISTINCT . . .**
- Force the result to be a bag (i.e., don't eliminate duplicates) by **ALL**, as in **. . . UNION ALL . . .**



Example: DISTINCT

- From Sells(bar, beer, price), find all the different prices charged for beers:

```
SELECT DISTINCT price  
FROM Sells;
```
- Notice that without DISTINCT, each price would be listed as many times as there were bar/beer pairs at that price.




Example: ALL

Force bag semantics with ALL

- Using relations `Frequents(drinker, bar)` and `Likes(drinker, beer)`:

```
(SELECT drinker FROM Frequents)  
EXCEPT ALL  
(SELECT drinker FROM Likes);
```
- Lists drinkers who frequent more bars than they like beers, and does so as many times as the difference of those counts.



Classroom Exercise

Assume R and S have the same schema

**Q1: (select * from R) INTERSECT ALL
(select * from S)**

**Q2: (select * from R) NATURAL JOIN (select
* from S)**

- a) Q1 and Q2 produce the same answer.
- b) The answer to Q1 is contained in the answer to Q2
- c) The answer to Q2 is contained in the answer to Q1
- d) Q1 and Q2 produce different answers.




Answer (b)

As sets, both produce the intersection of R and S . However, SQL is a bag language. Suppose a tuple t appears m times in R and n times in S . Then T appears $\min(m,n)$ times in the answer to Q1. However, for the join, we pair all tuples of R with all tuples of S that agree in the common attributes (i.e., all attributes in this case), and we produce of copy for each successful pairing. Thus, Q2 produces mn copies of t . It is easy to verify that as long as m and n are nonnegative integers, $\min(m,n) \leq mn$.




Join Expressions

- **SQL provides a number of expression forms that act like varieties of join in relational algebra.**
 - But using bag semantics, not set semantics.
- **These expressions can be stand-alone queries or used in place of relations in a FROM clause.**



Products and Natural Joins

- Natural join is obtained by:
R NATURAL JOIN S;
- Product is obtained by:
R CROSS JOIN S;
- Example:
Likes NATURAL JOIN Serves;
- Relations can be parenthesized subexpressions, as well.




Theta Join

- **$R \text{ JOIN } S \text{ ON } \langle \text{condition} \rangle$ is a theta-join, using $\langle \text{condition} \rangle$ for selection.**
- **Example: using Drinkers(name, addr) and Frequents(drinker, bar):**


**Drinkers JOIN Frequents ON
name = drinker;**

gives us all (d, a, d, b) quadruples such that drinker d lives at address a and frequents bar b .




Outerjoins

- **R OUTER JOIN S is the core of an outerjoin expression. It is modified by:**
 - 1. Optional NATURAL in front of OUTER.**
 - 2. Optional ON <condition> after JOIN.**
 - 3. Optional LEFT, RIGHT, or FULL before OUTER.**
 - ◆ **LEFT = pad dangling tuples of R only.**
 - ◆ **RIGHT = pad dangling tuples of S only.**
 - ◆ **FULL = pad both; this choice is the default.**



Database Modifications

- **A modification command does not return a result as a query does, but it changes the database in some way.**
- **There are three kinds of modifications:**
 - 1. *Insert* a tuple or tuples.**
 - 2. *Delete* a tuple or tuples.**
 - 3. *Update* the value(s) of an existing tuple or tuples.**



Insertion


- To insert a single tuple:
INSERT INTO <relation>
VALUES (<list of values>);
- Example: add to Likes(drinker, beer)
the fact that Sally likes Bud.

```
INSERT INTO Likes  
VALUES ( 'Sally' , 'Bud' );
```



Specifying Attributes in INSERT

- **We may add to the relation name a list of attributes.**
- **There are two reasons to do so:**
 1. **We forget the standard order of attributes for the relation.**
 2. **We don't have values for all attributes, and we want the system to fill in missing components with NULL or a default value.**



Example: Specifying Attributes

- **Another way to add the fact that Sally likes Bud to Likes(drinker, beer):**

```
INSERT INTO Likes (beer,  
drinker)  
VALUES ( 'Bud' ,  'Sally' ) ;
```



Inserting Many Tuples


- **We may insert the entire result of a query into a relation, using the form:**

**INSERT INTO <relation>
(<subquery>);**



Example: Insert a Subquery

- **Using `Frequents(drinker, bar)`, enter into the new relation `PotBuddies(name)` all of Sally's "potential buddies," i.e., those drinkers who frequent at least one bar that Sally also frequents.**



The other
drinker

Solution

Pairs of Drinker
tuples where the
first is for Sally,
the second is for
someone else,
and the bars are
the same.

INSERT INTO PotBuddies

(SELECT d2.drinker


FROM Frequents d1, Frequents d2

WHERE d1.drinker = 'Sally' AND

d2.drinker <> 'Sally' AND

d1.bar = d2.bar


);



Deletion

- To delete tuples satisfying a condition from some relation:

**DELETE FROM <relation>
WHERE <condition>;**



Example: Deletion

- Delete from Likes(drinker, beer) the fact that Sally likes Bud:

DELETE FROM Likes

**WHERE drinker = 'Sally' AND
beer = 'Bud' ;**



Example: Delete all Tuples

- **Make the relation Likes empty:**

```
DELETE FROM Likes;
```

- **Note no WHERE clause needed.**




Example: Delete Many Tuples

- **Delete from Beers(name, manf) all beers for which there is another beer by the same manufacturer.**

**DELETE FROM Beers b
WHERE EXISTS (**

**SELECT name FROM Beers
WHERE manf = b.manf AND
name <> b.name);**

Beers with the same manufacturer and a different name from the name of the beer represented by tuple b.




Semantics of Deletion -- 1

- **Suppose Anheuser-Busch makes only Bud and Bud Lite.**
- **Suppose we come to the tuple b for Bud first.**
- **The subquery is nonempty, because of the Bud Lite tuple, so we delete Bud.**
- **Now, When b is the tuple for Bud Lite, do we delete that tuple too?**



Semantics of Deletion -- 2

- **The answer is that we *do* delete Bud Lite as well.**
- **The reason is that deletion proceeds in two stages:**
 - 1. Mark all tuples for which the WHERE condition is satisfied in the original relation.**
 - 2. Delete the marked tuples.**



Updates


- To change certain attributes in certain tuples of a relation:

UPDATE <relation>

**SET <list of attribute
assignments>**

WHERE <condition on tuples>;

Updates many tuples at once.




Example: Update

- **Change drinker Fred's phone number to 555-1212:**

```
UPDATE Drinkers
```

```
SET phone = '555-1212'
```

```
WHERE name = 'Fred' ;
```




Example: Update Several Tuples

- **Make \$4 the maximum price for beer:**

UPDATE Sells

SET price = 4.00

WHERE price > 4.00;



Adding Attributes

- We may change a relation schema by adding a new attribute ("column") by:

```
ALTER TABLE <name> ADD  
    <attribute declaration>;
```

- Example:

```
ALTER TABLE Bars ADD  
phone CHAR(16) DEFAULT  
    'unlisted' ;
```



Deleting Attributes

- Remove an attribute from a relation schema by:

ALTER TABLE <name>

DROP <attribute>;

- Example: we don't really need the license attribute for bars:

ALTER TABLE Bars DROP license;

