

Chapter 6 The database


Language SQL –as a **tutorial**

- **About SQL**

SQL is a standard database language, adopted by many commercial systems.

ANSI SQL, SQL-92 or SQL2, **SQL99 or SQL3 extends SQL2 with object-relational features. SQL2003 is the collection of extensions to SQL3.**

- **How to query the database**
- **How to make modifications on database**
- **Transactions in SQL**



Subqueries

Result of a select-from-where query can be used in the where-clause of another query.


Simplest Case: Subquery Returns a Single, Unary Tuple

Find bars that serve Miller at the same price Joe charges for Bud.

Sells(bar, beer, price)

```
SELECT bar
FROM Sells
WHERE beer = 'Miller' AND price =
      (SELECT price
       FROM Sells
       WHERE bar = 'Joe's Bar' AND
         beer = 'Bud');
```

- Notice the **scoping rule**: an attribute refers to the most closely nested relation with that attribute.
- **Parentheses** around subquery are essential.



The IN Operator

“Tuple IN relation” is true iff the tuple is in the relation.

Example

Find the name and manufacturer of beers that Fred likes.

Beers(name, manf)

Likes(drinker, beer)

SELECT *

FROM Beers


WHERE name IN

(SELECT beer

FROM Likes

WHERE drinker = 'Fred');

● **Also: NOT IN.**



What is the difference?


```
SELECT a
FROM R, S
WHERE R.b = S.b;
```

We suppose:

R (a,b)

S (b,c)

```
SELECT a
FROM R
WHERE b IN (SELECT b FROM S);
```



IN is a Predicate About R's Tuples

```
SELECT a
FROM R
WHERE b IN (SELECT b FROM S);
```

Two 2's



One loop, over
the tuples of R


a	b
1	2
3	4

R

b	c
2	5
2	6

S

(1,2) satisfies
the condition;
1 is output once.



This Query Pairs Tuples from R, S

```
SELECT a
FROM R, S
WHERE R.b = S.b;
```

Double loop, over
the tuples of R and S


a	b
1	2
3	4

R

b	c
2	5
2	6


S

(1,2) with (2,5)
and (1,2) with
(2,6) both satisfy
the condition;
1 is output twice.



The Exists Operator

- **EXISTS(<relation>)** is true if and only if the <relation> is not empty.
- Being a Boolean-valued operator, **EXISTS** can appear in **WHERE** clauses.
- Example: From Beers(name, manf), find those beers that are the unique beer by their manufacturer.



Example Query with EXISTS

```
SELECT name  
FROM Beers b1  
WHERE NOT EXISTS(
```

Notice scope rule: manf refers to closest nested FROM with a relation having that attribute.

```
SELECT *  
FROM Beers  
WHERE manf = b1.manf AND  
      name <> b1.name);
```

Notice the SQL “not equals” operator

Set of beers with the same manf as b1, but not the same beer

● A subquery that refers to values from a surrounding query is called a **correlated subquery**.




The Operator ANY

- **$x = \text{ANY}(\langle \text{relation} \rangle)$ is a boolean condition meaning that x equals at least one tuple in the relation.**
- **Similarly, $=$ can be replaced by any of the comparison operators.**
- **Example: $x \geq \text{ANY}(\langle \text{relation} \rangle)$ means x is not smaller than all tuples in the relation.**
 - **Note tuples must have one component only.**



The Operator ALL

- Similarly, $x \neq \text{ALL}(\langle \text{relation} \rangle)$ is true if and only if for every tuple t in the relation, x is not equal to t .
 - That is, x is not a member of the relation.
- The \neq can be replaced by any comparison operator.
- Example: $x \geq \text{ALL}(\langle \text{relation} \rangle)$ means there is no tuple larger than x in the relation.



Quantifiers

ANY and ALL behave as existential and universal quantifiers, respectively.


Example

Find the beer(s) sold for the highest price.

Sells(bar, beer, price)


```
SELECT beer  
FROM Sells  
WHERE price >= ALL(  
    SELECT price  
    FROM Sells);
```

price from the outer
Sells must not be
less than any price.



Conditions Involving Relations

- **EXISTS R**: true if and only if R is not empty.
- **s IN R**: true if and only if s is equal to one of the values in R.
- **s > ALL R**: true if and only if s is greater than every value in unary R.
- **s > ANY R**: true if and only if s is greater than at least one value in unary R



Classroom exercise

Q1: select a from R


**Where $b \geq \text{ANY}$ (select d from S
where $c > 10$);**

Q2: select a from R

**Where $b \geq \text{ALL}$ (select d from S where
 $c > 10$);**

- a) Q1 and Q2 produce the same answer.
- b) The answer to Q1 is contained in the answer to Q2
- c) The answer to Q2 is contained in the answer to Q1
- d) Q1 and Q2 produce different answers.

Think about when the subquery is empty, what is the result?



Answer:

if the subquery is empty, then you can be bigger than ``all" yet there does not exist an element than which you are bigger. That's the way logic works.



Aggregations

Sum, avg, min, max, and count apply to attributes/columns. Also, count(*) applies to tuples.

- Use these in lists following SELECT.

Example

Find the average price of Bud.

Sells(bar, beer, price)

```
SELECT AVG(price)
```


```
FROM Sells
```

```
WHERE beer = 'Bud';
```

- Counts each tuple (presumably each bar that sells Bud) once.

Class Problem

What would we do if Sells were a bag?



Eliminating Duplicates Before Aggregation

Find the number of different prices at which Bud is sold.


Sells (bar, beer, price)

SELECT COUNT(DISTINCT price)

FROM Sells


WHERE beer = 'Bud';

- **DISTINCT may be used in any aggregation, but typically only makes sense with COUNT.**



NULL's Ignored in Aggregation

- **NULL never contributes to a sum, average, or count and can never be the minimum or maximum of a column.**
- **But if there are no non-NULL values in a column, then the result of the aggregation is NULL.**
- **Exception:** COUNT of an empty set is 0.



Examples

Select count(*)

From Sells

counts the number of tuples in Sells.

Select count(bar)


From Sells

**counts the number of values(non-NULL) in the bar column.
Duplicates values are not eliminated.**

Select count (distinct bar)

From Sells

**counts the number of different values in the bar column, no
matter how many kinds of beers bars sold.**




Example: Effect of NULL's

```
SELECT count(*)  
FROM Sells  
WHERE beer = 'Bud';
```

The number of bars
that sell Bud.

```
SELECT count(price)  
FROM Sells  
WHERE beer = 'Bud';
```

The number of bars
that sell Bud at a
known price.



Grouping

Follow select-from-where by GROUP BY and a list of attributes.

- **The relation that is the result of the FROM and WHERE clauses is grouped according to the values of these attributes, and aggregations take place only within a group.**

Example


Find the average sales price for each beer.

Sells(bar, beer, price)

SELECT beer, AVG(price)

FROM Sells

GROUP BY beer;



Example

Find, for each drinker, the average price of Bud at the bars they frequent.

Sells(bar, beer, price)

Frequents(drinker, bar)

SELECT drinker, AVG(price)

FROM Frequents, Sells


WHERE beer = 'Bud' AND

Frequents.bar = Sells.bar

GROUP BY drinker;

- Note: **grouping occurs after the \times and σ operations.**

Compute drinker-bar-price of Bud tuples first, then group by drinker



Restriction on SELECT Lists With Aggregation

- **If any aggregation is used, then each element of the SELECT list must be either:**
 - 1. Aggregated, or**
 - 2. An attribute on the GROUP BY list.**



Illegal Query Example


- You might think you could find the bar that sells Bud the cheapest by:

SELECT bar, MIN(price)

FROM Sells


WHERE beer = 'Bud';

- But this query is illegal in SQL.
 - Why? Note bar is neither aggregated nor on the GROUP BY list.




HAVING Clauses

- **HAVING <condition> may follow a GROUP BY clause.**
- **If so, the condition applies to each group, and groups not satisfying the condition are eliminated.**



Requirements on HAVING Conditions

- **These conditions may refer to any relation or tuple-variable in the FROM clause.**
- **They may refer to attributes of those relations, as long as the attribute makes sense within a group; i.e., it is either:**
 - 1. A grouping attribute, or**
 - 2. Aggregated.**



Example

Find the average price of those beers that are either served in at least 3 bars or manufactured by Anheuser-Busch.

Beers(name, manf)

Sells(bar, beer, price)

```
SELECT beer, AVG(price)
FROM Sells
GROUP BY beer
HAVING COUNT(*) >= 3 OR
       beer IN (
           SELECT name
           FROM Beers
           WHERE manf = 'Anheuser-Busch'
       );
```

Rules for having clause

- Anything goes in a subquery.
- Outside subqueries, they may refer to attributes only if they are either:

A grouping attribute, or
Aggregated




Grouping, Aggregation and Null

- The value NULL is ignored in any aggregation.
- NULL is treated as an ordinary value in a grouped attribute.

Select a, avg(b) from R

Result will be:

<u>a</u>	<u>avg(b)</u>
2	4
3	9
null	4



Home work

- Exercise 6.2.2 e)
- Exercise 6.3.1 c)
- Exercise 6.4.6 i)

***SELECT <desired attributes>**

***FROM <tuple variables or relation name>**

WHERE <conditions>

GROUP BY <attributes>

HAVING <conditions>

ORDER BY < list of attributes>

