

Chapter 6 The database


Language SQL –as a **tutorial**


- **About SQL**

SQL is a standard database language, adopted by many commercial systems.

ANSI SQL, SQL-92 or SQL2, **SQL99 or SQL3 extends SQL2 with object-relational features. SQL2003 is the collection of extensions to SQL3.**

- **How to query the database**
- **How to make modifications on database**
- **Transactions in SQL**





Transactions

What is transactions?

Why do we need transactions?

**How to set transaction with
different isolation level?**




Why Transactions?

- **Concurrent database access**

Execute *sequence of SQL statements* so they appear to be running in isolation

- **Resilience to system failures**

Guarantee all-or-nothing execution, regardless of failures




Concurrent Control

- **Database systems are normally being accessed by many users or processes at the same time.**
 - **Both queries and modifications.**


Serializability

Operations may be interleaved, but execution must be equivalent to *some* sequential (serial) order of all transactions



Resilience to system failures

- Failures may happen at any time.
- All or nothing done, never half done.



Transfer money from one account into another account.

Update accounts set balance = balance - 1000 where accounts.number=123;

Update accounts set balance = balance + 1000 where account.number= 456;




Solution for both concurrency and failures

Transactions

A transaction is a sequence of one or more SQL operations treated as a unit

- **Transactions appear to run in isolation**
- **If the system fails, each transaction's changes are reflected either entirely or not at all**



Transactions: SQL standard

- Normally with some strong properties regarding concurrency.
- Formed in SQL from single statements or explicit programmer control, i.e.


Transaction begins automatically on first SQL statement

- On “commit” transaction ends and new one begins.
- Current transaction ends on session termination.
- “Autocommit” turns each statement into transaction.




ACID Transactions

- ***ACID transactions*** are:
 - ***Atomic*** : Whole transaction or none is done.
 - ***Consistent*** : Database constraints preserved.
 - ***Isolated*** : It appears to the user as if only one process executes at a time.
 - ***Durable*** : Effects of a process survive a crash.





Consistency and isolation

- **Application defines consistency.**
- **Application requires isolation to achieve consistent results, there are four *isolation levels*.**
- **Locking typically used to achieve isolation.**




COMMIT

- The SQL statement COMMIT causes a transaction to complete.
 - It's database modifications are now permanent in the database.




ROLLBACK

- The SQL statement ROLLBACK also causes the transaction to end, but by *aborting*.
 - No effects on the database.
- Failures like division by 0 or a constraint violation can also cause rollback, even if the programmer does not request it.
- **Application** versus **system-generated** rollbacks.



Example: Interacting Processes

- Assume the usual **Sells(bar,beer,price)** relation, and suppose that Joe's Bar sells only Bud for \$2.50 and Miller for \$3.00.
- Sally is querying **Sells** for the highest and lowest price Joe charges.
- Joe decides to stop selling Bud and Miller, but to sell only Heineken at \$3.50.




Sally's Program

- Sally executes the following two SQL statements called **(min)** and **(max)** to help us remember what they do.

(max)SELECT MAX(price) FROM Sells
WHERE bar = 'Joe''s Bar';

(min)SELECT MIN(price) FROM Sells
WHERE bar = 'Joe''s Bar';



Joe's Program


- At about the same time, Joe executes the following steps: **(del)** and **(ins)**.

(del) DELETE FROM Sells

WHERE bar = 'Joe''s Bar';

(ins) INSERT INTO Sells

VALUES('Joe''s Bar', 'Heineken',
3.50);



Interleaving of Statements

- Although **(max)** must come before **(min)**, and **(del)** must come before **(ins)**, there are no other constraints on the order of these statements, unless we group Sally's and/or Joe's statements into transactions.



Example: Strange Interleaving

- Suppose the steps execute in the order **(max)(del)(ins)(min)**.

Joe's Prices:

{2.50,3.00} {2.50,3.00} {3.50}

Statement: (max) (del) (ins) (min)


Result: 3.00 3.50

- Sally sees MAX < MIN!




Fixing the Problem by Using Transactions

- If we group Sally's statements **(max)(min)** into one transaction, then she cannot see this inconsistency.
- She sees Joe's prices at some fixed time.
 - Either before or after he changes prices, or in the middle, but the MAX and MIN are computed from the same prices.




Another Problem: Rollback

- Suppose Joe executes **(del)(ins)**, not as a transaction, but after executing these statements, thinks better of it and issues a ROLLBACK statement.
- If Sally executes her statements after **(ins)** but before the rollback, she sees a value, 3.50, that never existed in the database.



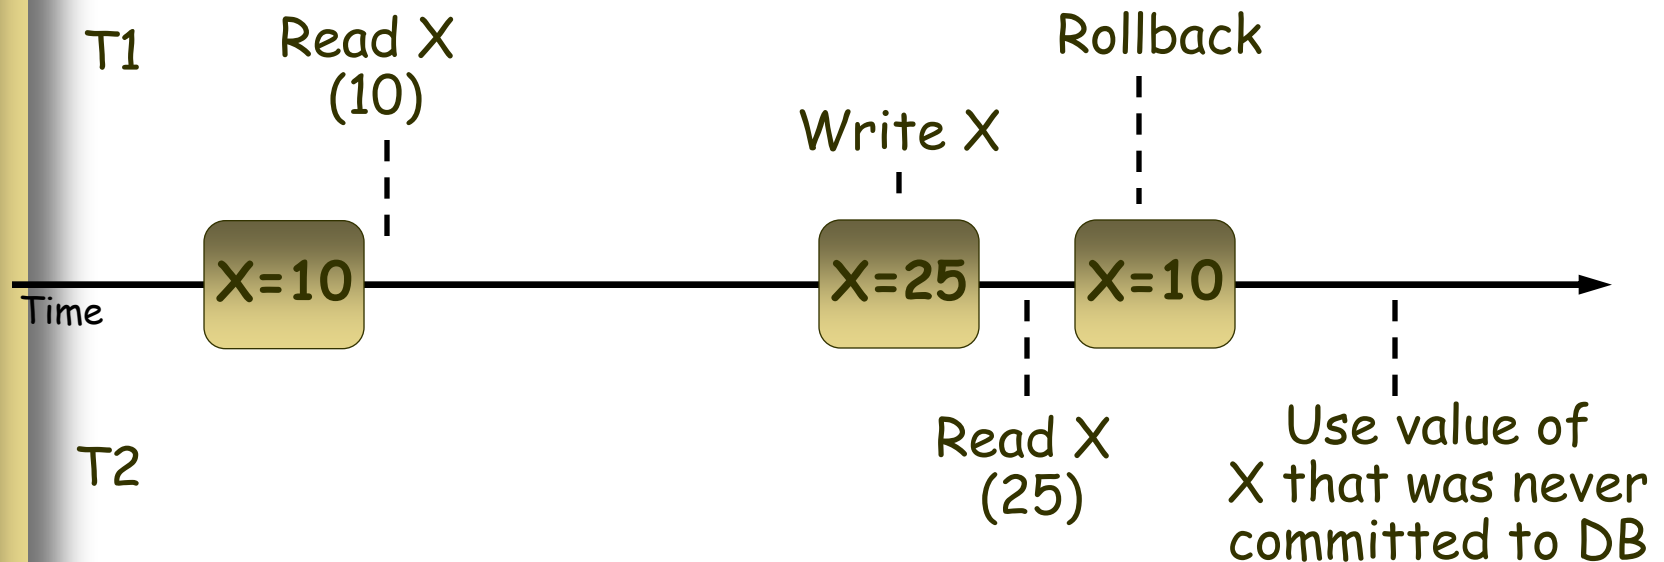
Solution

- If Joe executes **(del)(ins)** as a transaction, its effect cannot be seen by others until the transaction executes **COMMIT**.
 - If the transaction executes **ROLLBACK** instead, then its effects can *never* be seen.



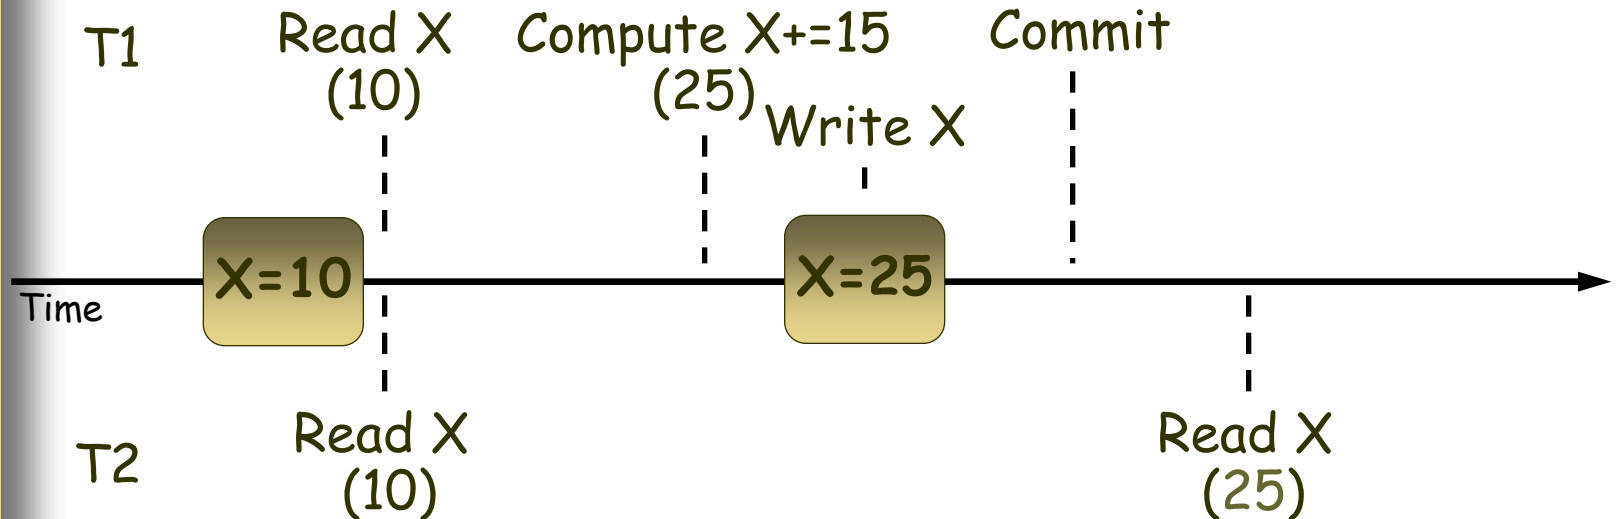
Summarize of problems caused by multiple users accessing (1)

• Dirty read




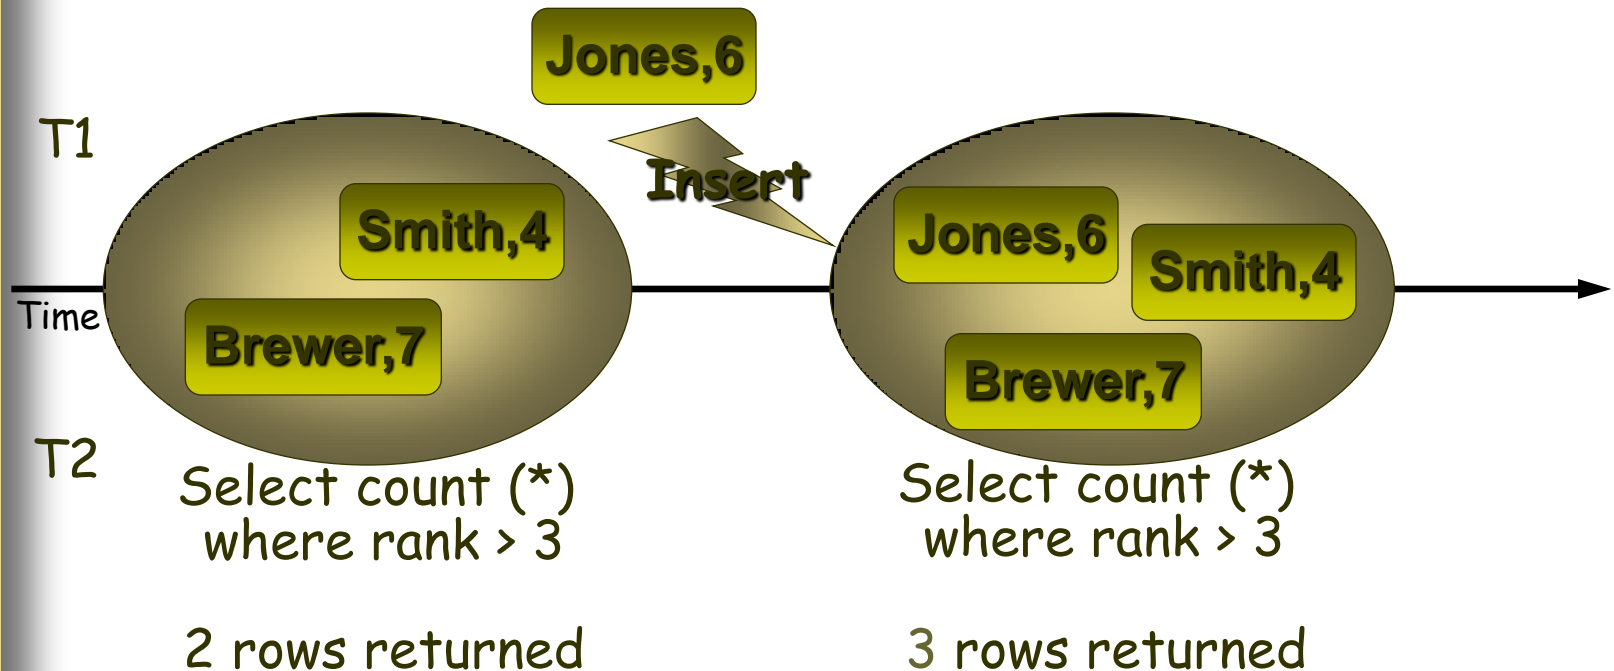
Summarize of problems caused by multiple users accessing (2)

- **Non-Repeatable Read**




Summarize of problems caused by multiple users accessing (3)

- The “Phantom” Problem



Isolation Levels

- SQL defines four *isolation levels* = choices about what interactions are allowed by transactions that execute at about the same time.
- Only one level ("serializable") = ACID transactions.
- Each DBMS implements transactions in its own way.



Choosing the Isolation Level

- Within a transaction, we can say:

**SET TRANSACTION ISOLATION
LEVEL X**

where $X =$

1. **SERIALIZABLE**

⇒ Overhead

⇒ Reduction in concurrency

2. **REPEATABLE READ**

3. **READ COMMITTED**

4. **READ UNCOMMITTED**


↓ Overhead ↑ Concurrency

↓ **Consistency Guarantees**




Serializable Transactions

- If Sally = **(max)(min)** and Joe = **(del)(ins)** are each transactions, and Sally runs with isolation level **SERIALIZABLE**, then she will see the database either before or after Joe runs, but not in the middle.




Isolation Level Is Personal Choice

- Your choice, e.g., run serializable, affects only how *you* see the database, not how others see it.
- **Example:** If Joe Runs serializable, but Sally doesn't, then Sally might see no prices for Joe's Bar.
 - i.e., it looks to Sally as if she ran in the middle of Joe's transaction.




Read-Committed Transactions

- If Sally runs with isolation level **READ COMMITTED**, then she can see only committed data, but not necessarily the same data each time.
- **Example**: Under **READ COMMITTED**, the interleaving **(max)(del)(ins)(min)** is allowed, as long as Joe commits.
 - Sally sees $MAX < MIN$.




Repeatable-Read Transactions

- Requirement is like read-committed, plus: if data is read again, then everything seen the first time will be seen the second time.
 - But the second and subsequent reads may see *more* tuples as well.



Example: Repeatable Read

- Suppose Sally runs under REPEATABLE READ, and the order of execution is (max)(del)(ins)(min).
 - (max) sees prices 2.50 and 3.00.
 - (min) can see 3.50, but must also see 2.50 and 3.00, because they were seen on the earlier read by (max).



Read Uncommitted

- A transaction running under **READ UNCOMMITTED** can see data in the database, even if it was written by a transaction that has not committed (and may never).
- **Example:** If Sally runs under **READ UNCOMMITTED**, she could see a price 3.50 even if Joe later aborts.



From weakest to strongest and the read behaviors they permit:


isolation level	dirty reads	nonrepeatable reads	phantoms
<hr/>			
READ UNCOMMITTED	Y	Y	Y
READ COMMITTED	N	Y	Y
REPEATABLE READ	N	N	Y
SERIALIZABLE	N	N	N

- True isolation is expensive in terms of concurrency
 - Many systems allow application to choose the phenomena they will live with
 - Trade off between correctness and concurrency




Homework

- **exercise 6.6.4**




Summary

- **SQL: The language is the principal query language for relational database systems. (SQL2, SQL3)**
- **Select-From-Where Queries**
- **Subqueries: The operators EXISTS, IN, ALL and ANY may be used to express boolean-valued conditions about the relations that are the result of a subquery**
- **Set Operations on Relations: UNION, INTERSECT, EXCEPT**



Summary(cont.)

- **The bag model for SQL, DISTINCT elimination of duplicate tuples; ALL allows the result to be a bag.**
- **Aggregations:
SUM,AVG,MIN,MAX,COUNT
GROUP BY, HAVING**
- **Modification Statements: INSERT,
DELETE, UPDATE**



SUMMARY(cont.)

- **Transactions: ACID**
- **Isolation levels :**
 1. **Serializable:** the transaction must appear to run either completely before or completely after each other transaction
 2. **Repeatable read:** every tuple read in response to a query will reappear if the query is repeated.
 3. **read-committed:** only tuples written by transactions that have already committed may be seen by the transaction.
 4. **Read-uncommitted:** no constraint.

