

Chapter 8 Views, Indexes

Virtual and Materialized Views
Speeding Accesses to Data

Views

- A *view* is a relation defined in terms of stored tables (called *base tables*) and other views.
- Two kinds:
 1. *Virtual* = not stored in the database; just a query for constructing the relation.
 2. *Materialized* = actually constructed and stored.

Declaring Views

- ❑ Declare by:

```
CREATE [MATERIALIZED] VIEW  
    <name> AS <query>;
```

- ❑ Default is virtual.

Example: View Definition

- `CanDrink(drinker, beer)` is a view “containing” the drinker-beer pairs such that the drinker frequents at least one bar that serves the beer:

```
CREATE VIEW CanDrink AS
  SELECT drinker, beer
  FROM Frequents, Sells
  WHERE Frequents.bar = Sells.bar;
```

Example: Accessing a View

- Query a view as if it were a base table.
 - Also: a limited ability to modify views if it makes sense as a modification of one underlying base table.
- Example query:

```
SELECT beer FROM CanDrink  
WHERE drinker = 'Sally';
```


What Happens When a View Is Used?

- **The DBMS starts by interpreting the query as if the view were a base table.**
 - Typical DBMS turns the query into something like relational algebra.
- **The queries defining any views used by the query are also replaced by their algebraic equivalents, and “spliced into” the expression tree for the query.**

Example: View Expansion

**SELECT beer
FROM CanDrink
WHERE drinker
= 'Sally';**

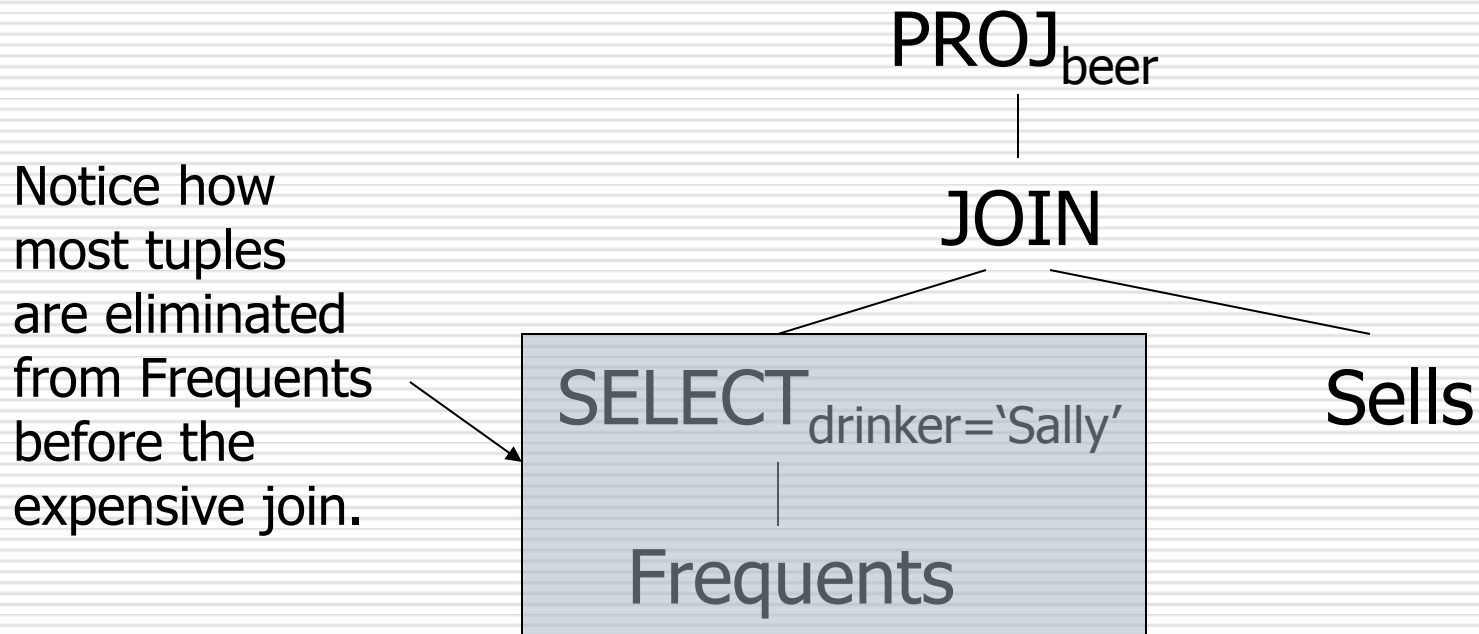
CREATE VIEW
CanDrink AS
SELECT drinker,
beer
FROM Frequents,
Sells
WHERE
Frequents.bar =
Sells.bar;



DMBS Optimization

- It is interesting to observe that the typical DBMS will then “optimize” the query by transforming the algebraic expression to one that can be executed faster.
- Key optimizations:
 1. Push selections down the tree.
 2. Eliminate unnecessary projections.

Example: Optimization



Modifying Views

☐ View Removal

Drop view canDrink;

- ☐ Updates on more complex views are difficult or impossible to translate, and hence are disallowed.
- ☐ Most SQL implementations allow updates only on simple views (without aggregates) defined on a single relation

Triggers on Views

- Generally, it is impossible to modify a virtual view, because it doesn't exist.
- But an INSTEAD OF trigger lets us interpret view modifications in a way that makes sense.
- **Example:** View Synergy has (drinker, beer, bar) triples such that the bar serves the beer, the drinker frequents the bar and likes the beer.

Example: The View

CREATE VIEW Synergy AS

SELECT Likes.drinker, Likes.beer, Sells.bar

Pick one copy of
each attribute



FROM Likes, Sells, Frequents

WHERE Likes.drinker =
Frequents.drinker

AND Likes.beer = Sells.beer

AND Sells.bar = Frequents.bar;



Natural join of Likes,
Sells, and Frequents

Interpreting a View Insertion

- ❑ We cannot insert into Synergy --- it is a virtual view.
- ❑ But we can use an INSTEAD OF trigger to turn a (drinker, beer, bar) triple into three insertions of projected pairs, one for each of Likes, Sells, and Frequents.
 - Sells.price will have to be NULL.

The Trigger

```
CREATE TRIGGER ViewTrig
  INSTEAD OF INSERT ON Synergy
  REFERENCING NEW ROW AS n
  FOR EACH ROW
  BEGIN
    INSERT INTO LIKES VALUES(n.drinker, n.beer);
    INSERT INTO SELLS(bar, beer) VALUES(n.bar, n.beer);
    INSERT INTO FREQUENTS VALUES(n.drinker, n.bar);
  END;
```

Materialized Views

- **Problem:** each time a base table changes, the materialized view may change.
 - Cannot afford to recompute the view with each change.
- **Solution:** Periodic reconstruction of the materialized view, which is otherwise “out of date.”

Example

```
CREATE MATERIALIZED VIEW CanDrink AS  
    SELECT drinker, beer  
    FROM Frequents, Sells  
    WHERE Frequents.bar = Sells.bar;
```


Indexes

- *Index* = data structure used to speed access to tuples of a relation, given values of one or more attributes.
- Could be a hash table, but in a DBMS it is always a balanced search tree with giant nodes (a full disk page) called a *B-tree*.

Declaring Indexes

- ❑ No standard!
- ❑ Typical syntax:

```
CREATE INDEX BeerInd ON  
  Beers (manf) ;
```

```
CREATE INDEX SellInd ON Sells (bar,  
  beer) ;
```

Using Indexes

- Given a value v , the index takes us to only those tuples that have v in the attribute(s) of the index.
- **Example**: use BeerInd and SellInd to find the prices of beers manufactured by Pete's and sold by Joe. (next slide)

Using Indexes --- (cont.)

```
SELECT price FROM Beers, Sells
WHERE manf = 'Pete''s' AND
      Beers.name = Sells.beer AND
      bar = 'Joe''s Bar';
```

1. Use BeerInd to get all the beers made by Pete's.
2. Then use SellInd to get prices of those beers, with bar = 'Joe''s Bar'

Database Tuning

- ❑ A major problem in making a database run fast is deciding which indexes to create.
- ❑ **Pro:** An index speeds up queries that can use it.
- ❑ **Con:** An index slows down all modifications on its relation because the index must be modified too.

Example: Tuning

- Suppose the only things we did with our beers database was:
 1. Insert new facts into a relation (10%).
 2. Find the price of a given beer at a given bar (90%).
- Then **SellInd** on Sells(bar, beer) would be wonderful, but **BeerInd** on Beers(manf) would be harmful.

Tuning Advisors

- A major research thrust.
 - Because hand tuning is so hard.
- An advisor gets a *query load*, e.g.:
 1. Choose random queries from the history of queries run on the database, or
 2. Designer provides a sample workload.

Tuning Advisors --- (2)

- The advisor generates candidate indexes and evaluates each on the workload.
 - Feed each sample query to the query optimizer, which assumes only this one index is available.
 - Measure the improvement/degradation in the average running time of the queries.

Some useful suggestions

- ❑ Index on its key.
- ❑ Index on the following two cases:
 1. If the attribute is almost a key
 2. If the tuples are **clustered** on that attribute.
- ➡ To decrease the cost of accessing data

Some useful suggestions (cont.)

- ❑ If we are doing mostly insertion, very few queries, then we do not want an index
- ❑ If we are doing a lot of queries, and the number of queries specifying movies and starts are roughly equally frequent, then both index are desired.

Summary of chapter 8

- ❑ Views (virtual and materialized)
- ❑ Updatable views
- ❑ Indexes (creation, use)